

10.1 Burrows-Wheeler transform

This exposition has been developed by David Weese. It is based on the following **sources**, which are all **recommended reading**:

1. M. Burrows and D. J. Wheeler (1994) *A Block-sorting Lossless Data Compression Algorithm*, SRC Research Report 124
2. G. Navarro and V. Mäkinen (2007) *Compressed Full-Text Indexes*, ACM Computing Surveys 39(1), Article no. 2 (61 pages), Section 5.3
3. D. Adjeroh, T. Bell, and A. Mukherjee (2008) *The Burrows-Wheeler Transform: Data Compression, Suffix Arrays and Pattern Matching*, Springer

10.2 Motivation

We have seen the suffix array to be an efficient data structure for exact string matching in a fixed text. However, for large texts like the human genome of about 3 billion basepairs, the text and the suffix array alone would consume **15 Gb** of memory. To solve the exact string matching in optimal $O(m + p)$ time (m =pattern length, p =number of occurrences) we would need an enhanced suffix array of $3 \cdot (1 + 4 + 4 + 4)$ Gb = **39 Gb** of memory.

Both cases exceed the amount of physical memory of a typical desktop computer, therefore we need a different data structure with a smaller memory footprint. Burrows and Wheeler proposed in 1994 a lossless compression algorithm (now used in bzip2). The algorithm transforms a given text into the so called *Burrows-Wheeler transform* (BWT), a permutation of the text that can be back transformed. The transformed text can in general be better compressed than the original text as in the BWT equal characters tend to form consecutive runs which can be compressed using run-length encoders.

We will see that it is possible to conduct exact searches using only the compressed BWT and some auxiliary tables.

10.3 Definitions

We consider a string T of length n . For $i, j \in \mathbb{N}$ we define:

- $[i..j] := \{i, i + 1, \dots, j\}$
- $[i..j) := [i..j - 1]$
- $T[i]$ is the i -th character of T .
- $T[i..j) := T[i]T[i + 1] \dots T[j)$ is the substring from the i -th to the j -th character
- We start counting from 1, i. e. $T = T[1..n]$
- $|T|$ denotes the string length, i. e. $|T| = n$
- The concatenation of strings X, Y is denoted as $X \cdot Y$, e. g. $T = T[1..i] \cdot T[i + 1..n]$ for $i \in [1..n)$

Definition 1 (cyclic shift). Let $T = T[1..n]$ be a text over the alphabet Σ that ends with unique character $T[n] = \$$, which is the lexicographically smallest character in Σ . The i -th *cyclic shift* of T is $T[i..n] \cdot T[1..i - 1]$ for $i \in [1..n]$ and denoted as $T^{(i)}$.

Example 2.

T	=	mississippi\$
$T^{(1)}$	=	mississippi\$
		\vdots
$T^{(3)}$	=	ssissippi\$mi
		\vdots
$T^{(n)}$	=	\$mississippi

10.4 Burrows-Wheeler transform

The Burrows-Wheeler transform (BWT) can be obtained by the following steps:

1. Form a conceptual matrix \mathcal{M} whose rows are the n cyclic shifts of the text T
2. Lexicographically sort the rows of \mathcal{M} .
3. Construct the transformed text T^{bwt} by taking the last column of \mathcal{M} .

The transformed text T^{bwt} in the last column is also denoted as L (last). Notice that every row and every column of \mathcal{M} , hence also the transformed text L is a permutation of T . In particular the first column of \mathcal{M} , call it F (first), is obtained by lexicographically sorting the characters of T (or, equally, the characters of L).

Example 3. Form \mathcal{M} and sort rows lexicographically:

		F	L
mississippi\$		\$	mississippi
ississippi\$m		i	\$mississipp
ssissippi\$mi		i	ppi\$mississ
sissippi\$mis		i	ssissippi\$mis
issippi\$miss		i	ssissippi\$m
ssippi\$missi	sort	m	ississippi\$
sippi\$missis	\Rightarrow	p	i\$mississip
ippi\$mississ		p	ppi\$mississi
ppi\$mississi		s	ippi\$missis
pi\$mississip		s	issippi\$mis
i\$mississipp		s	sippi\$missi
\$mississippi		s	ssissippi\$m

The transformed string L usually contains long runs of identical symbols and therefore can be efficiently compressed using move-to-front coding, in combination with statistical coders.

10.5 Constructing the BWT

Note that when we sort the rows of \mathcal{M} we are essentially sorting the suffixes of T . Hence, there is a strong relation between the matrix \mathcal{M} and the suffix array A of T .

i	$T^{(i)}$		i	$T^{(i)}$		A	$T[A[j]..n]$
1	mississippi\$		12	\$mississippi		12	\$
2	ississippi\$m		11	i\$mississipp		11	i\$
3	ssissippi\$mi		8	ippi\$mississ		8	ippi\$
4	sissippi\$mis		5	issippi\$miss		5	issippi\$
5	issippi\$miss		2	ississippi\$m		2	ississippi\$
6	ssippi\$missi	sort	1	mississippi\$	$\hat{=}$	1	mississippi\$
7	sippi\$missis	\Rightarrow	10	pi\$mississip		10	pi\$
8	ippi\$mississ		9	ppi\$mississi		9	ppi\$
9	ppi\$mississi		7	sippi\$missis		7	sippi\$
10	pi\$mississip		4	sissippi\$mis		4	sissippi\$
11	i\$mississipp		6	ssippi\$missi		6	ssippi\$
12	\$mississippi		3	ssissippi\$mi		3	ssissippi\$

Lemma 4. The Burrows-Wheeler transform T^{bwt} can be constructed from the suffix array A of T . It holds:

$$T^{\text{bwt}}[i] = \begin{cases} T[A[i]-1] & \text{if } A[i] > 1 \\ \$ & \text{else} \end{cases}$$

Proof: Since T is terminated with the special character $\$$, which is lexicographically smaller than any other character and occurs only at the end of T , a comparison of two shifts ends at latest after comparing a $\$$. Hence the characters right of the $\$$ do not influence the order of the cyclic shifts and they are sorted exactly like the suffixes of T . For each suffix starting at position $A[i]$ the last column contains the preceding character at position $A[i] - 1$ (or n resp.).

Corollary 5. The Burrows-Wheeler transform of a text of length n can be constructed in $O(n)$ time.

Corollary 6. The i -th row of \mathcal{M} contains the $A[i]$ -th cyclic shift of T , i. e. $\mathcal{M}_i = T^{(A[i])}$.

10.6 Reverse transform

One interesting property of the Burrows-Wheeler transform T^{bwt} is that the original text T can be reconstructed by a reverse transform without any extra information. Therefore we need the following definition:

Definition 7 (L-to-F mapping). Let \mathcal{M} be the sorted matrix of cyclic shifts of the text T . LF is a function $\text{LF} : [1..n] \rightarrow [1..n]$ that maps the rank of a cyclic shift X to the rank of $X^{(n)}$ which is X shifted by one to the right:

$$\text{LF}(l) = f \Leftrightarrow \mathcal{M}_f = \mathcal{M}_l^{(n)}$$

LF represents a one-to-one correspondence between elements of F and elements of L , and $L[i] = F[\text{LF}[i]]$ for all $i \in [1..n]$. Corresponding characters stem from the same position in the text. That can be concluded from the following equivalence:

$$\begin{aligned} \text{LF}(l) = f &\Leftrightarrow \mathcal{M}_f = \mathcal{M}_l^{(n)} \\ &\Leftrightarrow T^{(A[f])} = T^{(A[l])^{(n)}} \\ &\Leftrightarrow T^{(A[f])} = T^{(A[l] + n - 1)} \\ &\Leftrightarrow A[f] \equiv A[l] + (n - 1) \pmod{n} \end{aligned}$$

Example 8. $T = \text{mississippi\$}$. It holds $\text{LF}(1) = 2$ as the cyclic shift in row 1 of \mathcal{M} shifted by one to the right occurs in row 2. $\text{LF}(2) = 7$ as the cyclic shift in row 2 of \mathcal{M} shifted by one to the right occurs in row 7. For the same reason holds $\text{LF}(7) = 8$.

	F	L		i	LF(i)
$T^{(12)}$	\$ mississipp	i		1	2
$T^{(11)}$	i \$mississip	p		2	7
	i ppi\$missis	s		3	9
	i ssippi\$mis	s		4	10
	i ssissippi\$	m		5	6
	m \$ssissippi	\$		6	1
$T^{(10)}$	p i\$mississi	p		7	8
$T^{(9)}$	p p\$mississ	i		8	3
	s ippi\$missi	s		9	11
	s issippi\$mi	s		10	12
	s sippi\$miss	i		11	4
	s sissippi\$m	i		12	5

Thus the first character in row f stems from position $A[f]$ in the text. That is the same position the last character in row l stems from. One important observation is that the relative order of two cyclic shifts that end with the same character is preserved after shifting them one to the right.

Observation 9 (rank preservation). Let $i, j \in [1..n]$ with $L[i] = L[j]$. If $i < j$ then $\text{LF}[i] < \text{LF}[j]$ follows.

Proof: From $L[i] = L[j]$ and $i < j$ follows $\mathcal{M}_i[1..n-1] <_{\text{lex}} \mathcal{M}_j[1..n-1]$. Thus holds:

$$\begin{aligned} L[i] \cdot \mathcal{M}_i[1..n-1] &<_{\text{lex}} L[j] \cdot \mathcal{M}_j[1..n-1] \\ \Leftrightarrow \mathcal{M}_i^{(n)} &<_{\text{lex}} \mathcal{M}_j^{(n)} \\ \Leftrightarrow \mathcal{M}_{\text{LF}[i]} &<_{\text{lex}} \mathcal{M}_{\text{LF}[j]} \\ \Leftrightarrow \text{LF}[i] &< \text{LF}[j] \end{aligned}$$

Observation 9 allows to compute the LF-mapping without using the suffix array as the i -th occurrence of a character α in L is mapped to the i -th occurrence of α in F .

Example 10. $T = \text{mississippi\$}$. The L-to-F mapping preserves the relative order of indices of matrix rows that end with the same character. The increasing sequence of *all* indices $3 < 4 < 9 < 10$ of rows that end with s is mapped to the increasing and *contiguous* sequence $9 < 10 < 11 < 12$.

F	L	i	LF(i)
\$ mississipp	i	1	2
i \$mississip	p	2	7
i ppi\$missis	s	3	9
i ssippi\$mi	s	4	10
i ssissipp	m	5	6
m ississipp	\$	6	1
p i\$mississi	p	7	8
p i\$mississ	i	8	3
s ppi\$missi	s	9	11
s issippi\$mi	s	10	12
s slppi\$miss	i	11	4
s sissippi\$m	i	12	5

Definition 11. Let T be a text of length n over an alphabet Σ and L the BWT of T .

- Let $C : \Sigma \rightarrow [0..n]$ be a function that maps a character $c \in \Sigma$ to the total number of occurrences in T of the characters which are alphabetically smaller than c .
- Let $\text{Occ} : \Sigma \times [1..n] \rightarrow [1..n]$ be a function that maps $(c, k) \in \Sigma \times [1..n]$ to the number of occurrences of c in the prefix $L[1..k]$ of the transformed text L .

Theorem 12. For the L-to-F mapping LF of a text T holds:

$$LF(i) = C(L[i]) + \text{Occ}(L[i], i)$$

Proof: Let $\alpha = L[i]$. Of all occurrences of the character α in L , $\text{Occ}(L[i], i)$ gives index of the occurrence at position i starting counting from 1. $C(L[i]) + j$ is the position of the j -th occurrence of α in F starting counting from 1. With $j = \text{Occ}(L[i], i)$ the j -th occurrence of α in L is mapped to the j -th occurrence of α in F .

How can we back-transform L to T ? With the L-to-F mapping we reconstruct T from right to left by cyclic shifting by one to the right beginning with $T^{(n)}$ and extracting the first characters. That can be done using the following properties:

- The last character of T is $\$,$ whose only occurrence in F is $F[1]$, thus $\mathcal{M}_1 = T^{(n)}$ is T shifted by one to the right.
- $\mathcal{M}_{LF(1)} = T^{(n)} = T^{(n-1)}$ is T shifted by 2 to the right. Therefore $F[LF(1)]$ is the character before the last one in T .
- The character $T[n-i]$ is $F[\underbrace{LF(LF(\dots LF(1) \dots))}_{i \text{ times LF}}]$.

For the reverse transform we need F and the functions C and Occ . C and F can be obtained by bucket sorting L . LF only uses values $\text{Occ}(L[i], i)$ which can be precomputed in an array of size n by a sequential scan over L . The pseudo-code for the reverse transform of $L = T^{\text{bwt}}$ is given in algorithm **reverse transform**. We replaced F by L using $F[LF(i)] = L[i]$ and $F[1] = \$$.

Example 13. Reverse transform $L = \text{ipssm$piissii}$ of length $n = 12$ over the alphabet $\Sigma = \{\$, i, m, p, s\}$. First, we count the number of occurrences n_α of every character $\alpha \in \Sigma$ in L and compute the partial sums $C(\alpha) = \sum_{\beta < \alpha} n_\beta$ of characters smaller than α to obtain C .

$\alpha \in \Sigma$	\$	i	m	p	s
n_α	1	4	1	2	4
$C(\alpha)$	0	1	5	6	8

```

(1) // reverse_transform(L,Occ,C)
(2) i = 1, j = n, c = $;
(3) while (j > 0) do
(4)     T[j] = c;
(5)     c = L[i];
(6)     i = C(c) + Occ(c, i);
(7)     j = j - 1;
(8) od
(9) return T;

```

F is the concatenated sequence of runs of n_α many characters α in increasing order:

$$\begin{aligned}
 F &= \$^{n_s} \cdot i^{n_i} \cdot m^{n_m} \cdot p^{n_p} \cdot s^{n_s} \\
 &= \$iiiiimppssss
 \end{aligned}$$

For every i we precompute $\text{Occ}(L[i], i)$ by sequentially scanning L and counting the number of occurrences of $L[i]$ in L up to position i . That can be done during the first run, where we determine the values n_α .

i	1	2	3	4	5	6	7	8	9	10	11	12
$\text{Occ}(L[i], i)$	1	1	1	2	1	1	2	2	3	4	3	4

Begin in row $i = 1$:

Extract the character $T[n] = F[1] = \$$

$T = \dots\dots\dots \$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[1]) + \text{Occ}(L[1], 1) = 1 + 1 = 2$

Extract the character $T[n - 1] = F[2] = i$

$T = \dots\dots\dots i\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[2]) + \text{Occ}(L[2], 2) = 6 + 1 = 7$

Extract the character $T[n - 2] = F[7] = p$

$T = \dots\dots\dots pi\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[7]) + \text{Occ}(L[7], 7) = 6 + 2 = 8$

Extract the character $T[n - 3] = F[8] = p$

$T = \dots\dots\dots ppi\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[8]) + \text{Occ}(L[8], 8) = 1 + 2 = 3$

Extract the character $T[n - 4] = F[3] = i$

$T = \dots\dots\dots ippi\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[3]) + \text{Occ}(L[3], 3) = 8 + 1 = 9$

Extract the character $T[n - 5] = F[9] = s$

$T = \dots\dots\dots sippi\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[9]) + \text{Occ}(L[9], 9) = 8 + 3 = 11$

Extract the character $T[n - 6] = F[11] = s$

$T = \dots\dots\dots ssippi\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[11]) + \text{Occ}(L[11], 11) = 1 + 3 = 4$

Extract the character $T[n - 7] = F[4] = i$

$T = \dots\dots\dots issippi\$$

Proceed with row $i = \text{LF}(i)$:

$i = C(L[4]) + \text{Occ}(L[4], 4) = 8 + 2 = 10$

Extract the character $T[n - 8] = F[10] = s$

$T = \dots\dots\dots sissippi\$$

Proceed with row $i = \mathbf{LF}(i)$:

$$i = C(L[10]) + \text{Occ}(L[10], 10) = 8 + 4 = 12$$

Extract the character $T[n - 9] = F[12] = s$

$T = \dots\text{ssissippi}\$$

Proceed with row $i = \mathbf{LF}(i)$:

$$i = C(L[12]) + \text{Occ}(L[12], 12) = 1 + 4 = 5$$

Extract the character $T[n - 10] = F[5] = i$

$T = \dots\text{issippi}\$$

Proceed with row $i = \mathbf{LF}(i)$:

$$i = C(L[5]) + \text{Occ}(L[5], 5) = 5 + 1 = 6$$

Extract the character $T[1] = F[6] = m$

$T = \text{mississippi}\$$

10.7 Backward search

This exposition has been developed by David Weese. It is based on the following [sources](#), which are all recommended reading:

1. P. Ferragina, G. Manzini (2000) *Opportunistic data structures with applications*, Proceedings of the 41st IEEE Symposium on Foundations of Computer Science
2. P. Ferragina, G. Manzini (2001) *An experimental study of an opportunistic index*, Proceedings of the 12th ACM-SIAM Symposium on Discrete Algorithms, pp. 269-278

For a pattern $P = P[1..m]$ we want to count the number of occurrences in a text $T = T[1..m]$ given its Burrows-Wheeler transform $L = T^{\text{bwt}}$. If we would have access to the conceptual matrix \mathcal{M} we could conduct a binary search like in a suffix array. However, as we have direct access to only F and L we need a different approach.

Ferragina and Manzini proposed a backward search algorithm that searches the pattern from right to left by matching growing suffixes of P . It maintains an interval of matches and transforms the interval of matches of a suffix of P to an interval of the suffix which is one character longer. At the end, the length of the interval of the whole pattern P equals the number of occurrences of P .

Occurrences can be represented as intervals due to the following observation:

Observation 14. For every suffix of $P[j..m]$ of P the matrix rows \mathcal{M}_i with prefix $P[j..m]$ form a contiguous block. Thus there are $a, b \in [1..m]$ such that $i \in [a..b] \Leftrightarrow \mathcal{M}_i[1..m - j + 1] = P[j..m]$.

Proof: That is direct consequence of the lexicographical sorting of cyclic shifts in \mathcal{M} . Note that $[a..b] = \emptyset$ for $a > b$.

Consider $[a_j..b_j]$ to be the interval of matrix rows beginning with $P_j = P[j..m]$. In that interval we search cyclic shifts whose matching prefix is preceded by $c = P[j - 1]$ in the cyclic text, i. e. matrix rows that end with c . If we shift these rows by 1 to the right they begin with $P[j - 1..m]$ and determine the next interval $[a_{j-1}..b_{j-1}]$.

Matrix rows that end with c are occurrences of c in $L[a_j..b_j]$. The L-to-F mapping of the first and last occurrence yields a_{j-1} and b_{j-1} (rank preservation, lemma 9).

The L-to-F mapping maps the i -th c in L to the i -th c in F . How to determine i for the first and last c in $L[a_j..b_j]$ without scanning?

In L are $\text{Occ}(c, a_j - 1)$ occurrences of c before the first occurrence in $L[a_j..b_j]$, hence:

$$i_f = \text{Occ}(c, a_j - 1) + 1$$

i_l the number of the last occurrence of c in $L[a_j..b_j]$ equals the number of occurrences of c in $L[1..b_j]$:

$$i_l = \text{Occ}(c, b_j)$$

Now, we can determine a_{j-1} and b_{j-1} by:

$$\begin{aligned} a_{j-1} &= C(c) + \text{Occ}(c, a_j - 1) + 1 \\ b_{j-1} &= C(c) + \text{Occ}(c, b_j) \end{aligned}$$

```

(1) // count(P[1..m])
(2)  $i = m, a = 1, b = n;$ 
(3) while  $((a < b) \wedge (i \geq 1))$  do
(4)    $c = P[i];$ 
(5)    $a = C(c) + \text{Occ}(c, a - 1) + 1;$ 
(6)    $b = C(c) + \text{Occ}(c, b);$ 
(7)    $i = i - 1;$ 
(8) od
(9) if  $(b < a)$  then return "not found";
(10)   else return "found  $(b - a + 1)$  occurrences";
(11) fi

```

Algorithm **count** computes the number of occurrences of $P[1..m]$ in $T[1..n]$:

Example 15. $T = \text{mississippi\$}$. Search $P = \text{ssi}$.

Initialization: We begin with the empty suffix ϵ which is a prefix of every suffix, hence we initialize $a_{m+1} = 1$ and $b_{m+1} = n$.

	F	L
$a_4 \rightarrow$	\$ mississipp	i
	i \$mississip	p
	i ppi\$missis	s
	i ssippi\$mis	s
	i ssissippi\$	m
	m ississippi	\$
	p i\$mississi	p
	p pi\$mississ	i
	s ippi\$missi	s
	s issippi\$mi	s
	s sippi\$miss	i
$b_4 \rightarrow$	s sissippi\$m	i

$a_4 = 1$
 $b_4 = 12$

Searching P_3 : From all matrix rows we search those beginning with the last pattern character $P[3] = i$. From $\text{Occ}(x, 0) = 0$ and $\text{Occ}(x, n) = n_x$ follows $a_m = C(x) + 1$ and $b_m = C(x + 1)$.

	F	L		F	L
$a_4 \rightarrow$	\$ mississipp	i		\$ mississipp	i
	i \$mississip	p		i \$mississip	p
	i ppi\$missis	s		i ppi\$missis	s
	i ssippi\$mis	s		i ssippi\$mis	s
	i ssissippi\$	m		i ssissippi\$	m
	m ississippi	\$	\Rightarrow	m ississippi	\$
	p i\$mississi	p		p i\$mississi	p
	p pi\$mississ	i		p pi\$mississ	i
	s ippi\$missi	s		s ippi\$missi	s
	s issippi\$mi	s		s issippi\$mi	s
	s sippi\$miss	i		s sippi\$miss	i
$b_4 \rightarrow$	s sissippi\$m	i		s sissippi\$m	i

$a_4 = 1$
 $b_4 = 12$

$a_3 = C(i) + 1 = 1 + 1$
 $b_3 = C(i) + n_i = 1 + 4$

Searching P_2 : From all rows beginning with P_3 we search those beginning with $P[2] = s$. In L we count the s's in the part before the interval ($=0$) and including the interval ($=2$) to L-to-F map the first and last s in the interval.

Searching P_1 : From all matrix rows beginning with P_2 we search those beginning with $P[1] = s$.

Found the interval for P : $[a_1..b_1]$ is the interval of matrix rows with prefix P , thus P has $b_1 - a_1 + 1 = 2$ occurrences in the text T .

$a_1 \rightarrow$ **ss**issippi\$missi
 $b_1 \rightarrow$ **ssi**issippi\$mi
 $a_1 = 11$
 $b_1 = 12$

10.8 Locate matches

We have seen how to count occurrences of a pattern P in the text T , but how to obtain their location in the text? Algorithm **count** determines the indexes $a, a+1, \dots, b$ of matrix rows with prefix P . As cyclic shifts correspond to the suffixes of T , with a suffix array A of T we would be able to get the corresponding text position $pos(i)$ of the suffix in row i . It holds $pos(i) = A[i]$.

We will now see that it is not necessary to have given the whole suffix array of $4n$ bytes of memory. It suffices to have a fraction of the suffix array available to compute $pos(i)$ for every $i \in [1..n]$.

The idea is as following. We logically mark a suitable subset of rows in the matrix. For the marked rows we explicitly store the start positions of the suffixes in the text. If i is marked row $pos(i)$ is directly available. If i is not marked, the algorithm **locate** uses the L-to-F-mapping to find the row $i_1 = LF(i)$ corresponding to the suffix $T[pos(i) - 1..n]$. This procedure is iterated v times until we reach a marked row i_v for which $pos(i_v)$ is available; then we set $pos(i) = pos(i_v) + v$.

This is a direct space-time trade-off.

Example 16.

Example:

$pos(1) = 12$

$pos(3) = 8$

$pos(6) = 1$

$pos(10) = 4$

m	i	s	s	i	s	s	i	p	p	i	#
1	2	3	4	5	6	7	8	9	10	11	12

F
#
i
i
i
i
m
p
p
s
s
s

First
Last

Preprocessing. The position of every x -th letter in the text is marked and stored for the corresponding row in S .

For a row i , algorithm **locate** determines the location of the corresponding occurrence in $T[1..n]$.

```

(1) // locate(i)
(2)  $i' = i$ 
(3)  $v = 0$ ;
(4) while (row  $i'$  is not marked) do
(5)    $c = L[i']$ ;
(6)    $i' = C(c) + \text{Occ}(c, i')$ ;
(7)    $v = v + 1$ ;
(8) od
(9) return  $pos(i') + v$ ;

```

locate(i) is called for every $i = [a..b]$, where $[a..b]$ is interval of occurrences computed by **count**(P). We call the conjunction of both algorithms and their required data structures the *FM Index*.

In the following we give an example using the BWT of the text `mississippi#`, the thinned out suffix array and the interval $[a..b]$ resulting from the search of the pattern `si`.

For each $i = [9, 10]$ we have to its position in the text. **i = 9**

Step 1

F		L	
#	mississipp	i	1
i	#mississip	p	2
i	ppi#missis	s	3
i	ssippi#mis	s	4
i	ssissippi#	m	5
m	ississippi	#	6
p	i#mississi	p	7
p	pi#mississ	i	8
s	ippi#missi	s	9
s	issippi#mi	s	10
s	sippi#miss	i	11
s	sissippi#m	i	12

row 9 is not marked
 $\rightarrow L\text{-to-}F(9)=11$
 \rightarrow Look at row 11
 $v=1$

Step 2

F		L	
#	mississipp	i	1
i	#mississip	p	2
i	ppi#missis	s	3
i	ssippi#mis	s	4
i	ssissippi#	m	5
m	ississippi	#	6
p	i#mississi	p	7
p	pi#mississ	i	8
s	ippi#missi	s	9
s	issippi#mi	s	10
s	sippi#miss	i	11
s	sissippi#m	i	12

row 11 is not marked
 $\rightarrow L\text{-to-}F(11)=4$
 \rightarrow Look at row 4
 $v=2$

Step 3

F		L	
#	mississipp	i	1
i	#mississip	p	2
i	ppi#missis	s	3
i	ssippi#mis	s	4
i	ssissippi#	m	5
m	ississippi	#	6
p	i#mississi	p	7
p	pi#mississ	i	8
s	ippi#missi	s	9
s	issippi#mi	s	10
s	sippi#miss	i	11
s	sissippi#m	i	12

row 4 is not marked
 $\rightarrow L\text{-to-}F(4)=10$
 \rightarrow Look at row 10
 $v=3$

row 10 is marked
 Calculation of the $pos(9)$:
 $pos(9) = pos(10) + 3 = 4 + 3 = 7$

We saw how to avoid storing the complete suffix array when locating the text.

However, the table *Occ* is still quite big. It contains the number of occurrences for each character and each possible prefix of L needing $|\Sigma| \times n$ entries storing the number of occurrences.

One way to reduce the size of *Occ* is to store only every x -th index. The entries that are omitted can be reconstructed from stored entries at the cost of an increased running time, by simply counting in the BWT from the last stored position on.

Taken together, we can conduct an exact search in a text in time *linear* to the query size.

For example (for DNA) using the text T (n bytes), the BWT (n bytes resp. $n/4$ bytes), the *Occ* table (e.g. $4 \cdot 4 \cdot n/32$ bytes when storing only every 32th entry) and a sampled suffix array (e.g. $4 \cdot n/8$ bytes when marking every 8th entry). In our example calculation we would need about $2.25n - 3n$ bytes.

11.9 Compressing the FM Index

This exposition has been developed by David Weese. It is based on the following sources, which are all recommended reading:

1. P. Ferragina, G. Manzini (2000) *Opportunistic data structures with applications*, Proceedings of the 41st IEEE Symposium on Foundations of Computer Science
2. P. Ferragina, G. Manzini (2001) *An experimental study of an opportunistic index*, Proceedings of the 12th ACM-SIAM Symposium on Discrete Algorithms, pp. 269-278
3. Johannes Fischer (2010), *Skriptum VL Text-Indexierung*, SoSe 2010, KIT
4. A. Andersson (1996) *Sorting and searching revisited*, Proceedings of the 5th Scandinavian Workshop on Algorithm Theory, pp. 185-197

11.10 RAM Model

From now on we assume the *RAM model* in which we model a computer with a CPU that has registers of w bits which can be modified with logical and arithmetical operations in $O(1)$ time. The CPU can directly access a memory of at most 2^w words.

In the following we assume $n \leq 2^w$ so that it is possible to address the whole input. To have a more precise measure, we count memory consumptions in bits. The uncompressed suffix array then does not require $O(n)$ memory but $O(n \log n)$ bits, as $\lceil \log_2 n \rceil$ bits are required to represent any number in $[1..n]$.

11.11 Tables of the FM Index

Let T be a text of length n over the alphabet Σ and $\sigma = |\Sigma|$ be the alphabet size. We have seen, that for the algorithms **count** and **locate** we need L and the tables C and Occ . Without compression their memory consumption is as follows:

- $L = T^{\text{bwt}}$ is a string of length n over Σ and requires $O(n \log \sigma)$ bits
- C is an array of length σ over $[0..n]$ and requires $O(\sigma \log n)$ bits
- Occ is an array of length $\sigma \times n$ over $[0..n]$ and requires $O(\sigma \cdot n \log n)$ bits
- pos (if every row is marked) is a suffix array of length n over $[1..n]$ and requires $O(n \log n)$ bits

We will present approaches to compress L , Occ and pos , but omit to compress C assuming that σ and $\log n$ are tolerably small.

11.12 Compressing L

Burrows and Wheeler proposed a move-to-front coding in combination with Huffman or arithmetic coding. In the context of the move-to-front encoding each character is encoded by its index in a list, which changes over the course of the algorithm. It works as follows:

1. Initialize a list Y of characters to contain each character in Σ exactly once
2. Scan L with $i = 1, \dots, n$
 - (a) Set $R[i]$ to the number of characters preceding character $L[i]$ in the list Y
 - (b) Move character $L[i]$ to the front of Y

R is the MTF encoding of L . R can again be decoded to L in a similar way (Exercise).

Algorithm **move_to_front(L)** shows the pseudo-code of the move-to-front encoding. The array M maintains for every alphabet character the number preceding characters in Y instead of using Y directly.

```

(1) // move_to_front(L)
(2) for j = 1 to  $\sigma$  do
(3)    $M[j] = j - 1$ 
(4) od
(5) for i = 1 to n do
(6)   // ord maps a character to its rank in the alphabet
(7)    $x = \text{ord}(L[i])$ 
(8)    $R[i] = M[x]$ ;
(9)   for j = 1 to  $\sigma$  do
(10)    if  $M[j] < M[x]$  then  $M[j] = M[j] + 1$ ; fi
(11)  od
(12)   $M[x] = 0$ ;
(13) od
(14) return R;

```

Observation 17. The BWT tends to group characters together so that the probability of finding a character close to another instance of the same character is increased substantially:


final char (L)	sorted rotations
a	n to decompress. It achieves compression
o	n to perform only comparisons to a depth
o	n transformation} This section describes
o	n transformation} We use the example and
o	n treats the right-hand side as the most
a	n tree for each 16 kbyte input block, enc
a	n tree in the output stream, then encodes
i	n turn, set $L[i]$ to be the
i	n turn, set $R[i]$ to the
o	n unusual data. Like the algorithm of Man
a	n use a single set of probabilities table
e	n using the positions of the suffixes in
i	n value at a given point in the vector R
e	n we present modifications that improve t
e	n when the block size is quite large. Ho
i	n which codes that have not been seen in
i	n with $\$ch\$$ appear in the {\em same order
i	n with $\$ch\$$. In our exam
o	n with Huffman or arithmetic coding. Bri
o	n with figures given by Bell\cite{bell}.

Observation 18. The move-to-front encoding replaces equal characters that in L are “close together” by “small values” in R . In practice, the most important effect is that zeroes tend to occur in runs in R . These can be compressed using an order-0 compressor, e.g. the Huffman encoding.

i	$L[i]$	$R[i]$	Y_{next}
			aeio
1	a	0	aeio
2	o	3	oaei
3	o	0	oaei
4	o	0	oaei
5	o	0	oaei
6	a	1	aoei
7	a	0	aoei
8	i	3	iaoe
9	i	0	iaoe
10	o	2	oiae
11	a	2	aoie
12	e	3	eaio
13	i	3	ieao
14	e	1	eiao
15	e	0	eiao
16	i	1	ieao
17	i	0	ieao

...

The Huffman encoding builds a binary tree where leaves are alphabet characters. The tree is balanced such that for every node the leaves in the left and right subtree have a similar sum of occurrences.



Left and right child are labeled with 0 and 1. The labels on the paths to each leaf define its bit code. The more frequent a character the shorter its bit code. The final sequence H is the bitwise concatenation of bit codes of characters from left to right in R .

The final sequence of bits H is:

$$\begin{aligned}
 L &= \text{a o o o a a i i} \dots \\
 R &= \text{0 3 0 0 1 0 3 0} \dots \\
 H &= \text{0 1 0 0 0 1 1 0 0 1 0 0} \dots
 \end{aligned}$$

One property of the MTF coding is that the whole prefix $R[1..i-1]$ is required to decode character $R[i]$, the same holds for H . For encoding and decoding this is fine (practical assignment).

However, we want to search in the compressed FM index and hence need random accesses to L in algorithm **locate** which would take $O(n)$ time. Manzini and Ferragina achieve this directly on the Huffman encoded R , however their algorithm is not practical, albeit optimal in theory.

We will proceed differently by using a simple trick we can determine $L[i]$ using the Occ function. Clearly, the values $\text{Occ}(c, i)$ and $\text{Occ}(c, i-1)$ differ only for $c = L[i]$.

Thus we can determine both $L[i]$ and $\text{Occ}(L[i], i)$ using σ Occ-queries, which we will see take in sum $O(\sigma)$ time. Using wavelet trees this time can even be reduced to $O(\log \sigma)$.

11.13 Compressing Occ

We reduce the problem of counting the occurrences of a character in a prefix of L to counting 1's in a prefix of a bitvector. Therefore we construct a bitvector B_c of length n for each $c \in \Sigma$ such that:


$$B_c[i] = \begin{cases} 1 & \text{if } L[i] = c \\ 0 & \text{else} \end{cases}.$$

Definition 19. For a bitvector B we define $\text{rank}_1(B, i)$ to be the number of 1's in the prefix $B[1..i]$. $\text{rank}_0(B, i)$ is defined analogously.

As each 1 in the bitvector B_c indicates an occurrence of c in L , it holds:


$$\text{Occ}(c, i) = \text{rank}_1(B_c, i).$$

We will see that it is possible to answer a rank query of a bitvector of length n in constant time using additional tables of $o(n)$ bits. Hence the σ bitvectors are an implementation of Occ that allows to answer Occ queries in constant time with an overall memory consumption of $O(\sigma n + o(\sigma n))$ bits. Given a bitvector $B = B[1..n]$. We compute the length $\ell = \lfloor \frac{\log n}{2} \rfloor$ and divide B into blocks of length ℓ and superblocks of length ℓ^2 .



presents a bitvector with some real values. The superblocks are on top.

The following picture



1. For the i -th superblock we count the number of 1's from the beginning of B to the end of the superblock in $M'[i] = \text{rank}_1(B, i \cdot \ell^2)$. As there are $\lfloor \frac{n}{\ell^2} \rfloor$ superblocks, M' can be stored in $O(\frac{n}{\ell^2} \cdot \log n) = O(\frac{n \log n}{\ell^2}) = o(n)$ bits.
2. For the i -th block we count the number of 1's from the beginning of the overlapping superblock to the end of the block in $M[i] = \text{rank}_1(B[1 + k\ell..n], (i - k)\ell)$ where $k = \lfloor \frac{i-1}{\ell} \rfloor$ is the number of blocks left of the overlapping superblock. M has $\lfloor \frac{n}{\ell} \rfloor$ entries and can be stored in $O(\frac{n}{\ell} \cdot \log \ell^2) = O(\frac{n \log \log n}{\log n}) = o(n)$ bits.
3. Let P be a precomputed lookup table such that for each possible bitvector V of length ℓ and $i \in [1..\ell]$ holds $P[V][i] = \text{rank}_1(V, i)$. V has $2^\ell \times \ell$ entries of values at most ℓ and thus can be stored in

$$O(2^\ell \cdot \ell \cdot \log \ell) = O\left(2^{\frac{\log n}{2}} \cdot \log n \cdot \log \log n\right) = O(\sqrt{n} \log n \log \log n) = o(n)$$

bits.

We now decompose a rank-query into 3 subqueries using the precomputed tables. For a position i we determine the index $p = \lfloor \frac{i-1}{\ell} \rfloor$ of next block left of i and the index $q = \lfloor \frac{p-1}{\ell} \rfloor$ of the next superblock left of block p . Then it holds:

$$\text{rank}_1(B, i) = M'[q] + M[p] + P[B[1 + p\ell..(p+1)\ell]][i - p\ell] .$$

Note that $B[1 + p\ell..(p+1)\ell]$ fits into a single CPU register and can therefore be determined in $O(1)$ time. Thus a rank-query can be answered in $O(1)$ time.

11.14 Compressing pos

To compress pos we mark only a subset of rows in the matrix M and store their text positions. Therefore we need a data structure that efficiently decides whether a row $M_i = T[j]$ is marked and that retrieves j for a marked row i .

If we would mark every η -th row in the matrix ($\eta > 1$) we could easily decide whether row i is marked, e. g. iff $i \equiv 1 \pmod{\eta}$. Unfortunately this approach still has worst-cases where a single pos -query takes $O(\frac{\eta-1}{\eta}n)$ time (exercise).

Instead we mark the matrix row for every η -th text position, i. e. for all $j \in [0..\lceil \frac{n}{\eta} \rceil]$ row i with $M_i = T^{(1+j\eta)}$ is marked with the text position $pos(i) = 1 + j\eta$. To determine whether a row is marked we could store all marked pairs $(i, 1 + j\eta)$ in a hash map or a binary search tree with key i .

11.15 Compressing pos

Again we can use a $O(1)$ rank-query supported bitvector B_{pos} in conjunction with an array Pos of size n/η .

If we still have the suffix array during the construction of the BWT, we can simply scan through the array maintaining an index k which initialize to 0. Whenever $A[i]/\eta \equiv 0 \pmod{n}$, we mark the i -th Bit in the Bitvector, store $Pos[k] = A[i]$ and increment k .

During a second scan of the suffix array, whenever $A[i]/\eta \equiv 0 \pmod{n}$ we look up how many bits k were set before i by using the bitvector. Then we set $Pos[k] = A[i]$ and hence sparsify the suffix array.

If the suffix array is not given, we use the BWT and L-to-F mapping traverse the BWT as in the reconstruction algorithm. While doing this, we keep counting the number of steps. After η backwards steps, we are at textposition $n - \eta$ and hence we mark the bitvector.

After setting all bits we again traverse the BWT maintaining a counter m which we initialize with 0. Whenever the bitvector is set we increment m , obtain the rank k of the bit and set $Pos[k] = n - m \cdot \eta$

11.16 Appendix: Packed B-tree

Ferragina and Manzini proposed a different approach. They marked every η -th text position for $\eta = \Theta(\log^2 n)$ and divided the matrix in buckets of η adjacent rows. For each marked row they recorded the row offset to the first row of the bucket. This offset takes $O(\log \eta) = O(\log \log n)$ bits.

As each bucket has at most η marked rows they use a *packed B-tree* (Appendix) of $u = O(\log^2 n)$ keys of size $k = O(\log \log n)$ bits. This B-tree supports membership queries in $O(\log_{w/k} u) = O\left(\frac{\log \log n}{\log \log n - \log \log \log n}\right) = O(1)$ time.

Each packet B-tree uses space proportional to the number of stored keys. Hence the *pos* data structure has an overall space consumption of $O\left(\frac{n}{\eta}(\log \log n + \log n)\right)$ bits since with each marked row M_i they also keep the value $\text{pos}(i)$ using $O(\log n)$ bits. A *packed B-tree* (Andersson 1996) is a balanced search tree whose nodes store keys of k bits length. Inner nodes store $t = \lfloor \frac{w}{k+1} \rfloor$ keys $y_1 < y_2 < \dots < y_t$ in sorted order and have $t + 1$ children. It is easy to see that searching a node in a tree of u leaves then takes $O(\log t \cdot \log_{t+1} u)$ as the tree height is $O(\log_{t+1} u)$ and at each level, the search chooses the child pointer (subtree) whose separation values are on either side of the search value in $O(\log t)$ (binary search on node keys).

The main trick of the packed B-tree is to replace the binary search of the child pointer by a bit-parallel comparison of all the keys in constant time. Let the sorted keys $y_1 < \dots < y_t$ of a node be stored in a register Y ascending from left to right each separated by a 1 bit. Let X be a register with a similar layout storing t copies of the search key x each separated by a 0 bit. For the difference $Y - X$ holds that the bit at the separating position left of node key y_i is 0 iff $y_i < x$. If we clear all but the separating bits by logically AND'ing $(Y - X)$ with a corresponding mask M , the result contains a sequence of r 0 bits followed by $t - r$ 1 bits, where r is the rank of x among the keys y_1, \dots, y_t .

$$\begin{array}{l}
 Y \quad 1|00010|1|00111|1|01001|1|01110|1|10101|1|11000|1|11011|1|11110| \\
 X \quad 0|01011|0|01011|0|01011|0|01011|0|01011|0|01011|0|01011|0|01011| \\
 Y - X \quad 0|10111|0|11100|0|11110|1|00011|1|01010|1|01101|1|10000|0|10011| \\
 M \quad 1|00000|1|00000|1|00000|1|00000|1|00000|1|00000|1|00000|1|00000| \\
 (Y - X) \text{ AND } M \quad 0|00000|0|00000|0|00000|1|00000|1|00000|1|00000|1|00000|1|00000|
 \end{array}$$

Lemma 20. *The rank of x among the keys of a node can be determined in $O(1)$.*

Proof: M can be precomputed. X can be computed from the search key x by $X = x * (M \gg t)$. We compute $(Y - X) \text{ AND } M$ in $O(1)$. The rank of x can then be determined in $O(1)$ by looking up the result in a precomputed lookup table.

Hence, the whole packed B-tree search takes $O(\log_{t+1} u) = O(\log_{w/k} u)$ time.